



v1 — cross-attention into a 256-d ring buffer (2 of 36 layers) — failed at deployment



*Only 2 of 36 layers see the bridge; no LLM inductive bias for 256-d vectors.*

v2 — LLaVA-style prepend into the 4096-d input space (all 36 layers) — works



*Bridge tokens live in the LLM's input embedding space; full causal attention from all layers.*