


v1 — Cross-attention into 256-d ring buffer (failed: KL=0.004 offline, $L=225 < F=256$ at deployment)



v2 — LLaVA-style token-prepend (works: $L > T$ by 26–82% on 4 games)

